

Multi-objective Bayesian optimisation for design of Pareto-optimal current drive profiles in STEP

Theodore Brown^{1,2}, Stephen Marsden¹, Francis Casson¹, Alexander Terenin² and Vignesh Gopakumar¹

¹United Kingdom Atomic Energy Authority, Oxford, UK OX14 3EB
²Department of Engineering, University of Cambridge, UK CB2 1PZ

Introduction

JETTO is a high-fidelity plasma modelling code that solves coupled core transport and equilibrium equations. It is used to **evaluate plasma scenarios** as part of the STEP design process. JETTO is initialised with the parameters of a candidate design, and the simulation is run until the plasma reaches a steady state. The steady-state properties of the plasma are used to assess the **impact of design choices** and the **suitability** of a given design.

However, JETTO takes **several hours** to run, which severely limits the extent to which the design space can be explored.

We demonstrate **multi-objective Bayesian optimisation**, a method for design optimisation that delivers **higher-quality solutions** in significantly **fewer iterations** than previous methods used in STEP, using techniques from **machine learning** and **surrogate modelling**. Our approach also offers **improved interpretability**, allowing design engineers to **quantify the tradeoffs** between different objectives.

Our example application is the optimisation of electron cyclotron resonance heating profiles to achieve desirable safety factor properties.




Figure 1: Overview

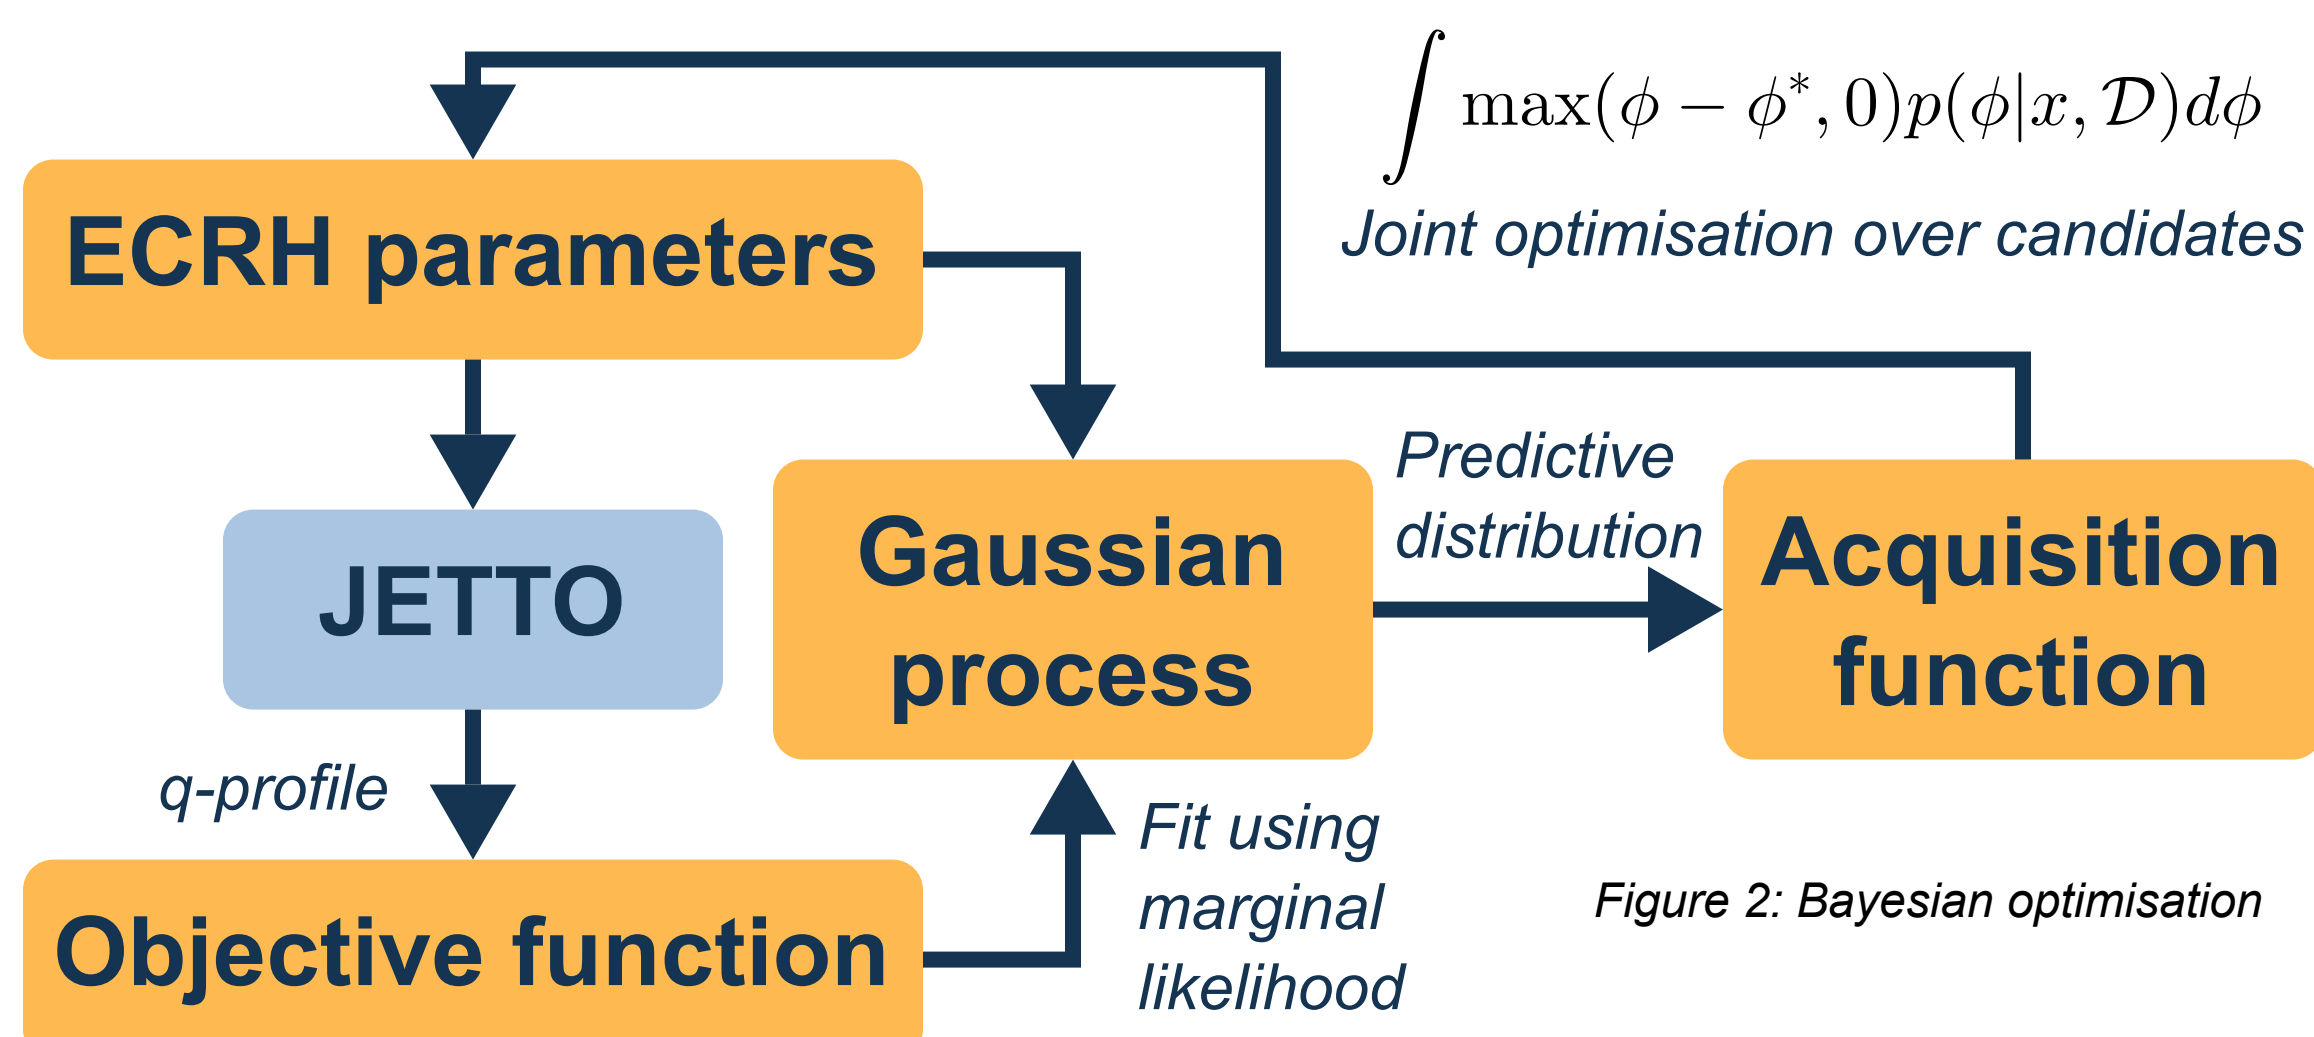


Figure 2: Bayesian optimisation

What is Bayesian optimisation?

Bayesian optimisation is built upon a **probabilistic model**, such as a **Gaussian Process (GP)**. The model is fit to the **previous observations** (inputs and objective values). It is used to generate **predictions** about the **performance of unseen points**. An inner optimisation loop uses the predictions to find the **most promising points** to try next.

This method:

- reduces** the number of runs 'wasted' on trying suboptimal points
- ensures that all the **information** gained at each step is **propagated** to future steps

As a result, BO exhibits **vastly improved performance** compared to stochastic search methods.




Figure 3: Using a GP model to select next candidates


What is Pareto optimality?

Single-objective optimisation tasks involve finding **one solution** that maximises a **scalar objective**. Comparing solutions is easy, as each solution is either better or worse than the others.

In **multi-objective** settings, there is normally no solution that simultaneously maximises every objective. Instead, we seek to find a **set of solutions** that represent the **tradeoffs** between each objective. The set is made up of all points that are **Pareto optimal**.

A point is **Pareto optimal** if it is impossible to improve its performance under one objective without reducing its performance under another objective.

Figure 4: Pareto optimality criterion



Target metrics for safety factor profile

Property	Rationale	Formulation
Minimise shear at centre	Improve β	$q(0) - \min q \quad \arg \min_{\rho} q$
Minimum q above 2	Improve β	$\ \min q - (2 + \epsilon) \ $
Monotonic q	Improve β	$\int_0^1 q \mathbf{1} \left(\frac{dq}{d\rho} \leq 0 \right) d\rho$
Monotonic gradient of q	Reduce fast ion losses	$\int_0^1 \frac{dq}{d\rho} \mathbf{1} \left(\frac{d^2q}{d\rho^2} \leq 0 \right) d\rho$
High shear at integer q	Mitigate NTMs	$q^{-1}(n), n \in \mathbb{N}$

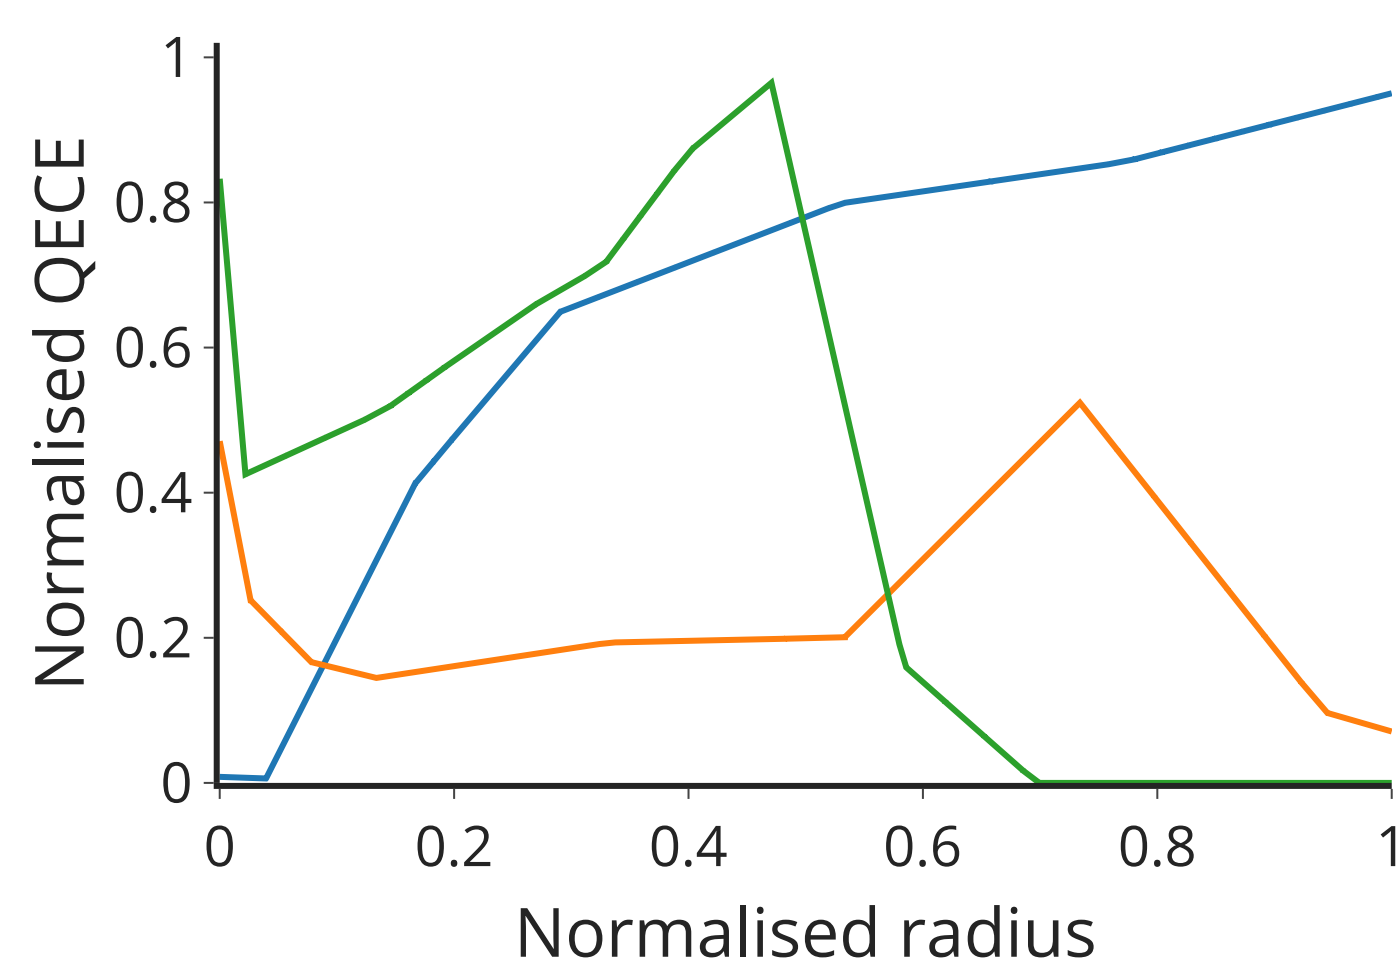


Figure 5: Example piecewise linear ECRH profiles

Parameterisation of ECRH profile

The choice of input parameterisation affects the **rate of convergence** and **quality of results**. Ideally, the parameterisation would be **general**, so that every possible profile can be represented. However, this can mean that the search space is **too large** to explore effectively.

Previously, the ECRH profile has been represented as a piecewise linear function (Fig. 5), with manually tuned constraints on the parameters. We also experimented with using a **sum-of-Gaussians representation** (Fig. 6), where each Gaussian represents an EC beam launcher.

The SoG parameterisation produces smooth ECRH profiles, and leads to a simpler mapping from input space to objective space: **small changes in parameters result in small changes in objective value**. This ensures that the mapping can be well-represented by a GP model.

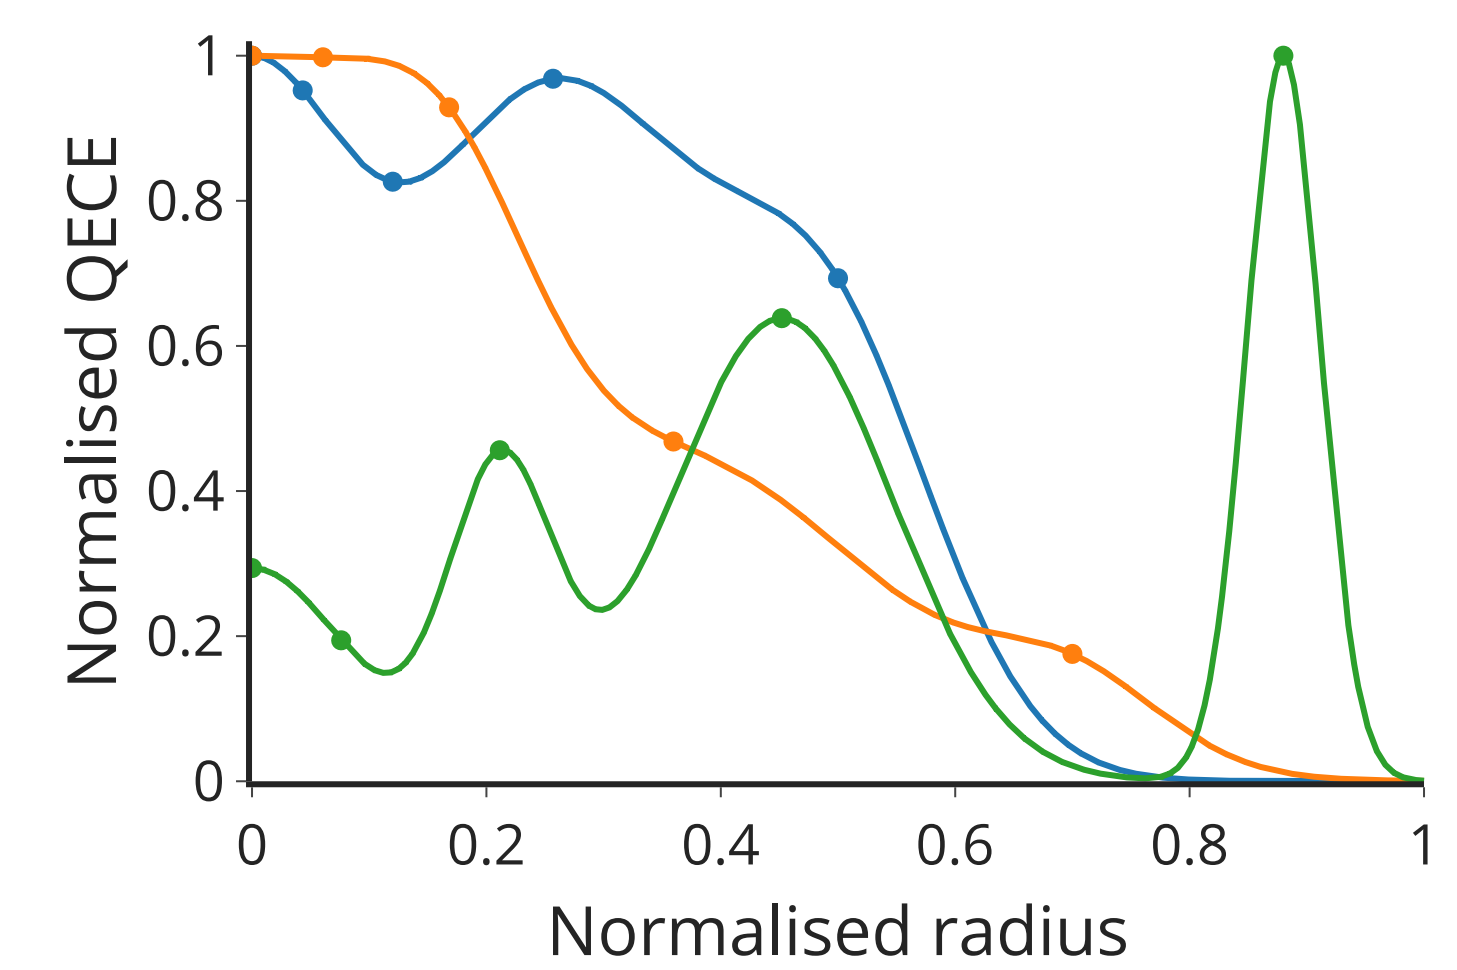


Figure 6: Example sum-of-Gaussians ECRH profiles

Results with piecewise linear ECRH

Running the multi-objective Bayesian optimisation loop using the piecewise linear ECRH function for 5 steps with 5 candidates at each step produces 8 Pareto optimal solutions.

The solutions are split into two groups: **monotonically decreasing ECRH profiles** (e.g. Solution 1) and **double-peaked ECRH profiles** (e.g. Solution 8). STEP designs currently use monotonically decreasing ECRH, as previous optimisation methods failed to find good double-peaked solutions.



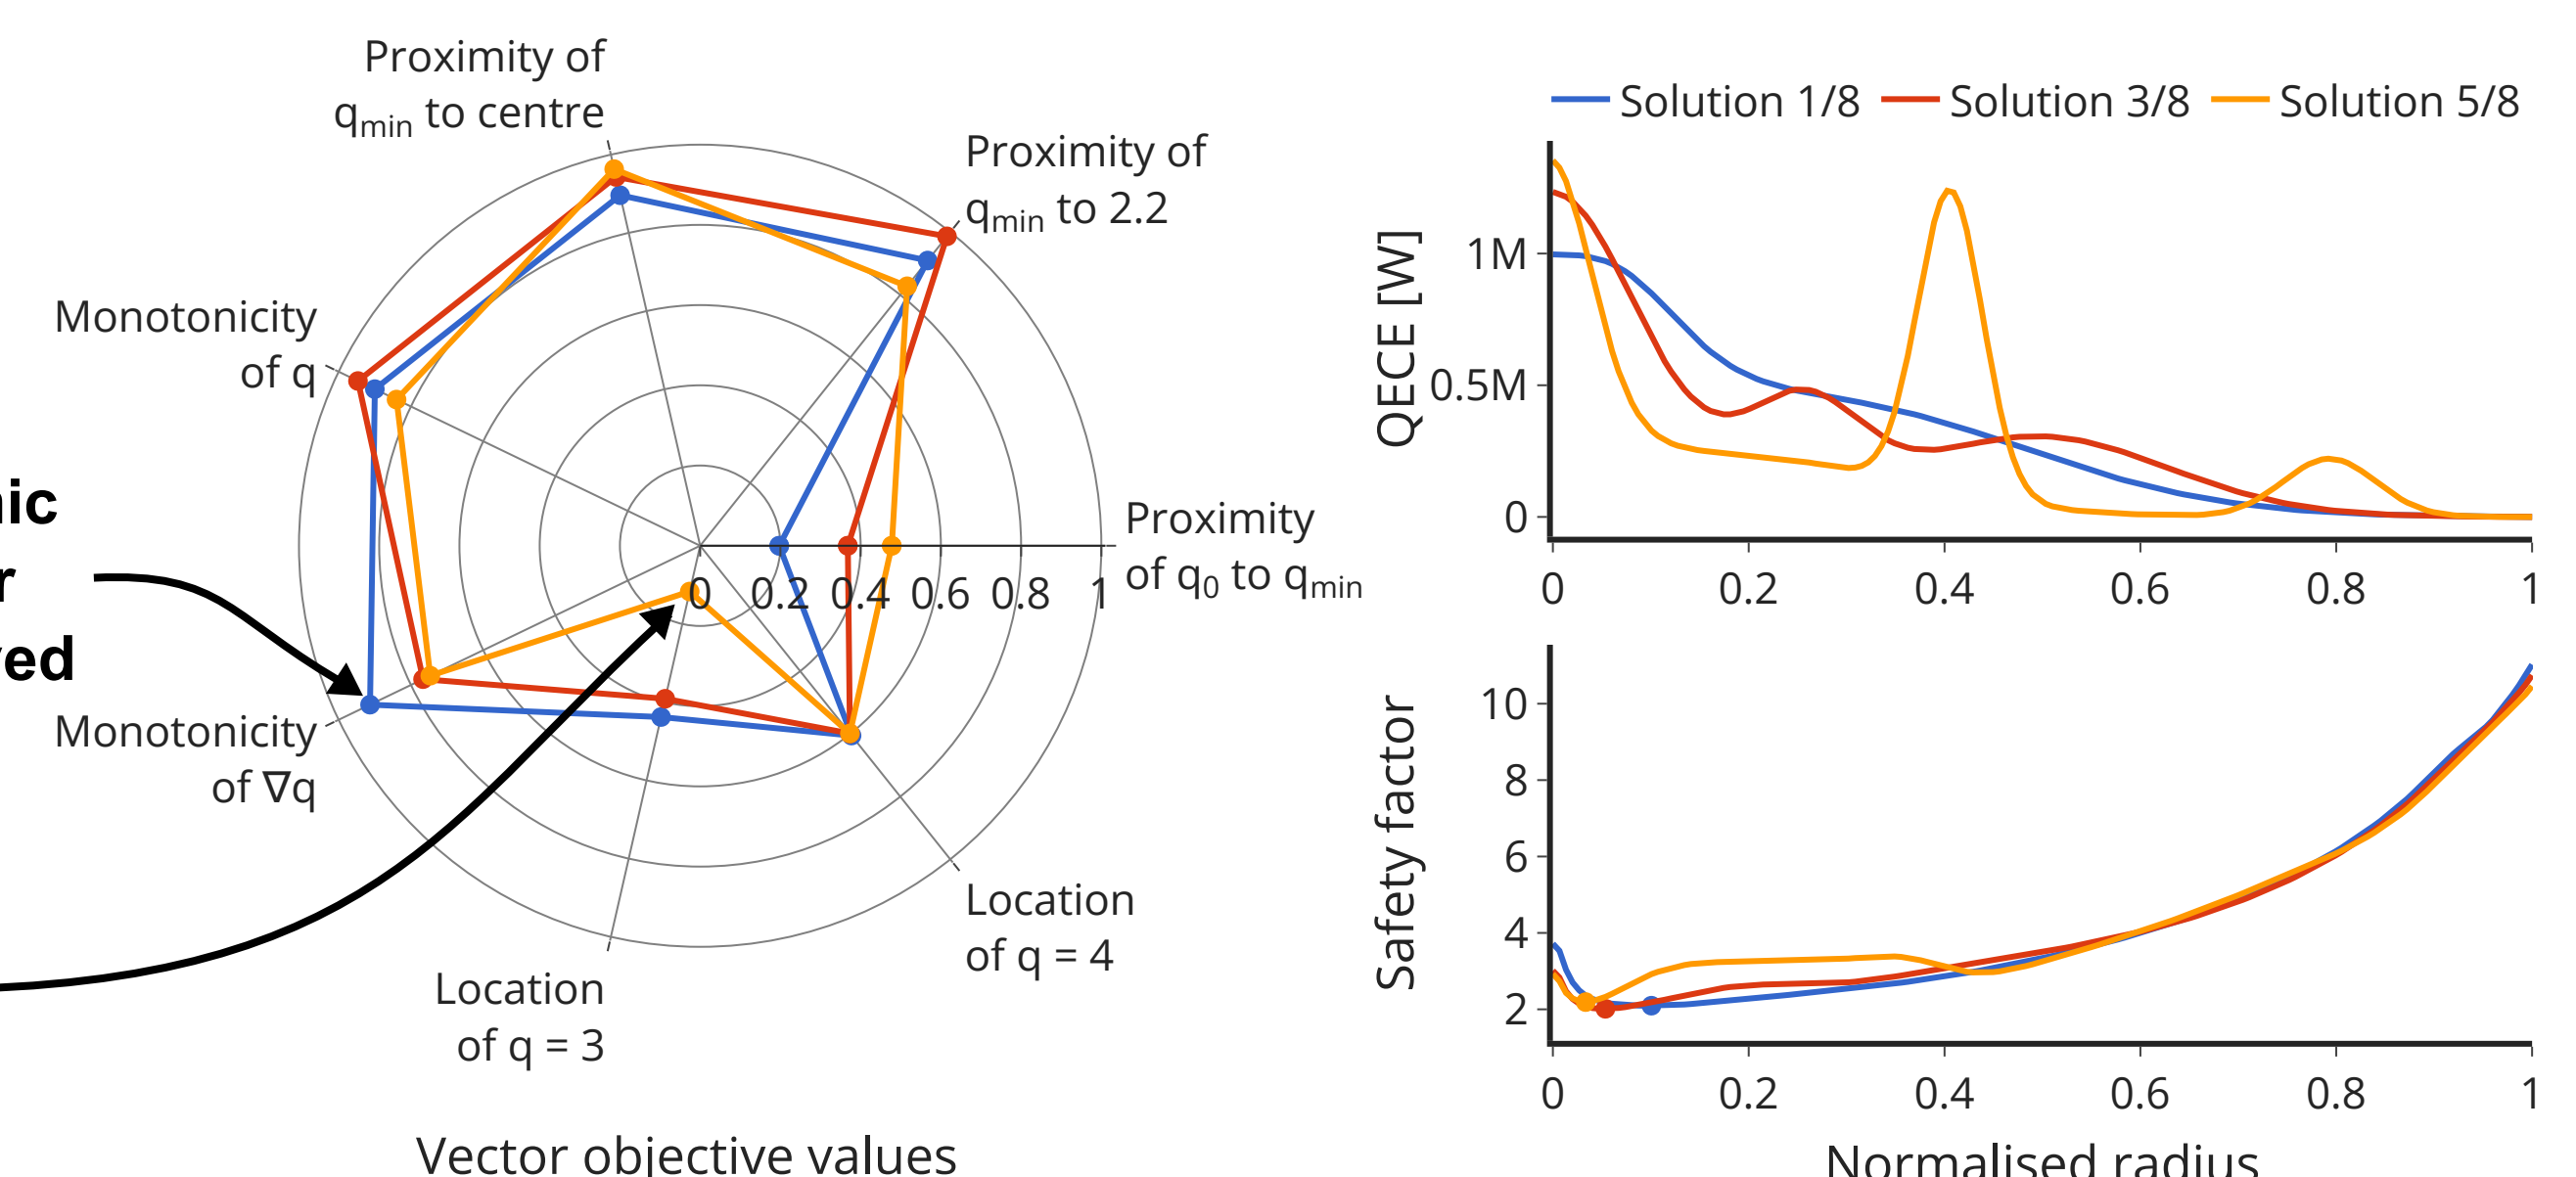


Figure 7: A selection of Pareto-optimal solutions generated by MOBO with piecewise linear ECRH (L) and sum-of-Gaussians ECRH (R)

Results with sum-of-Gaussians ECRH

Running the multi-objective Bayesian optimisation loop using the sum-of-Gaussians ECRH function for 10 steps with 30 candidates at each step produces 8 Pareto optimal solutions. **More steps** are required because the parameterisation is **more general**.

The ECRH parameterisation we used has up to 4 Gaussians logarithmically-spaced in $[0, 1]$. Achieving the **monotonicity objectives** are much more challenging with this representation.



STEP OPEN

This project was funded by the UKAEA/STFC FARSCAPE programme

The UK Atomic Energy Authority's mission is to lead the commercial development of fusion power and related technology, and position the UK as a leader in sustainable nuclear energy

Find out more
www.gov.uk/ukaea

Contact: tab3@cam.ac.uk

Future work

- Better objective and ECRH functions
- MOBO analysis of **ballooning stability**
- Trained models for **transfer learning**